

spaceTimeWindow Library

Technical Manual

Version 1.0 (Draft) | OpenFOAM v2512

Boundary Data Extraction and Flow Recalculation • **Time Interpolation** • **Integrated Zstd Compression** • **Optional Encryption**

Alin-Adrian Anton

Politehnica University of Timișoara

Last updated: February 6, 2026

OPENFOAM® is a registered trademark of OpenCFD Limited, producer and distributor of the OpenFOAM software via <https://www.openfoam.com>.

This work is licensed under CC BY-SA 4.0.
<https://creativecommons.org/licenses/by-sa/4.0/>

Copyright © 2026 Alin-Adrian Anton

This work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International License (CC BY-SA 4.0).

You are free to:

- **Share** – copy and redistribute the material in any medium or format
- **Adapt** – remix, transform, and build upon the material for any purpose, even commercially

Under the following terms:

- **Attribution** – You must give appropriate credit, provide a link to the license, and indicate if changes were made.
- **ShareAlike** – If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.

<https://creativecommons.org/licenses/by-sa/4.0/>

OPENFOAM® is a registered trademark of OpenCFD Limited.

CHERUBIM® is a registered trademark of the author. This software is libre software (libreware).

Contents

1	Introduction	6
1.1	Requirements	6
1.2	Building the Library	6
2	Concept: Space-Time Window Extraction and Reconstruction	7
2.1	The Challenge of Large-Scale CFD Simulations	7
2.2	The Space-Time Window Concept	7
2.3	Extraction and Reconstruction Workflow	8
2.3.1	Extraction Phase (HPC)	8
2.3.2	Reconstruction Phase (Offline)	9
2.4	Parallel Execution with MPI Domain Decomposition	9
2.4.1	How Parallel Extraction Works	9
2.5	Data Storage: Initial Fields and Boundary Data	11
2.5.1	Initial Fields	11
2.5.2	Boundary Data	12
2.6	Why Reconstruction Works	12
2.7	Automatic Boundary Condition Configuration	12
2.7.1	Resolution-Independent Reconstruction	13
2.7.2	Time Interpolation at Runtime	13
3	Encryption of Boundary Data	15
3.1	Key Generation	15
3.2	Secure Key Storage	15
3.3	Encrypted Extraction	16
3.4	Decryption During Case Initialization	16
3.5	Security Properties	16
3.5.1	Cryptographic Foundations	17
3.6	Ethical Considerations and HPC Transparency	17
3.7	Protection Against Malware and Ransomware	18
3.7.1	Secure Archival Strategy	18
3.7.2	Security Model	19
4	Boundary Condition Approaches	19
4.1	Pressure-Coupled BC — Recommended	19
4.2	Inlet-Outlet BC (-inletOutletBC) — For Quick Tests	20
4.3	Fixed Outlet Direction (-outletDirection)	21
4.4	Boundary Condition Assignment Logic	21
5	Workflow Overview	21
5.1	Serial Workflow	22
5.2	Parallel Workflow	22
6	Configuration Reference	23
6.1	Extraction Configuration	23
6.2	Initial Fields vs Boundary Data Fields	24
6.3	Boundary Condition Configuration	25
6.3.1	spaceTimeWindowCoupledPressure (Recommended for Pressure)	26
6.3.2	spaceTimeWindowInletOutlet	27
6.3.3	spaceTimeWindow (Pure Dirichlet)	28

6.4	Case Initialization Options	29
6.5	Mesh Coarsening and Refinement	30
6.5.1	Spatial Interpolation Algorithms	30
6.5.2	Initial Field Interpolation	32
7	Data Storage Format	34
7.1	Directory Structure	34
7.2	Boundary Data Compression	34
7.2.1	Delta-Varint Codec (DVZ)	35
7.2.2	DVZ Mathematical Formulation	36
7.2.3	Delta-Varint-Temporal Codec (DVZT)	37
7.2.4	DVZT Mathematical Formulation	37
7.2.5	Integrated Zstd Compression	40
7.2.6	External Compression Benchmark	41
7.3	Extraction Metadata	42
8	Parallel Execution	43
8.1	Parallel Extraction	43
8.2	Field Reconstruction	43
8.3	Parallel Reconstruction	44
9	Mass Conservation	44
9.1	The Mass Imbalance Problem	44
9.2	Solutions	44
9.2.1	Use Pressure-Coupled BC (Recommended)	44
9.2.2	Use <code>-inletOutletBC</code> (for quick tests)	45
9.2.3	Use <code>-correctMassFlux</code>	45
9.2.4	Use <code>-outletDirection</code> with <code>fixesValue true</code>	45
9.3	The <code>fixesValue</code> Option	45
9.4	Outlet Patch for Pressure Relief	46
9.5	Recommended Configurations	46
10	Solver Settings and Relaxation	46
10.1	Time Stepping	46
10.2	PIMPLE Settings	47
10.3	Relaxation Factors	47
10.4	Linear Solver Settings	48
11	Time Interpolation	48
11.1	Exact Timestep Matching (Default)	50
11.2	Linear Interpolation	50
11.3	Cubic Spline Interpolation	51
11.3.1	Catmull-Rom Spline Formulation	51
11.4	Time Range Requirements	52
12	Flux Reporting and Diagnostics	52
13	Troubleshooting	52
13.1	Common Errors	52
13.1.1	Time Step Mismatch	52
13.1.2	No Matching Timestep	53
13.1.3	Pressure Solver Divergence	53

13.1.4	Encryption Errors	53
14	Acknowledgments	53
A	Example Case: ERCOFTAC UFR2-02 Square Cylinder	54
A.1	Reference Data and Validation Resources	54
A.2	Physical Background: Von Kármán Vortex Street	54
A.3	Domain and Mesh Configuration	55
A.4	Extraction Box Placement	55
A.5	Physical Parameters	56
A.6	Running the Example	56
A.7	Validation	56
B	Field Selection by Turbulence Model	57

1 Introduction

The `spaceTimeWindow` library enables extraction and reconstruction of spatial subsets from Large Eddy Simulation (LES) and transient RANS simulations in OpenFOAM. This technique allows researchers to:

- Extract a region of interest from a full-domain simulation
- Store time-varying boundary conditions efficiently with compact codecs (DVZ/DVZT)
- Optionally apply integrated zstd compression for additional $\sim 10\times$ size reduction
- Reconstruct the flow within the subset using pre-computed boundary data
- Optionally encrypt boundary data for secure distribution

The library consists of four main components:

1. **`spaceTimeWindowExtract`** – Function object for boundary data extraction during simulation
2. **`spaceTimeWindowInitCase`** – Utility for reconstruction case initialization
3. **`spaceTimeWindow`** – Pure Dirichlet boundary condition for applying extracted data
4. **`spaceTimeWindowInletOutlet`** – Flux-based boundary condition (recommended for unsteady flows)

1.1 Requirements

- OpenFOAM v2512 (openfoam.com) or compatible ESI-OpenCFD version
- C++17 compatible compiler (GCC 7+ or Clang 5+)
- Optional: libzstd for integrated zstd compression (libzstd-dev or libzstd-devel)
- Optional: libsodium for encryption support (libsodium-dev or libsodium-devel)
- Optional: MPI for parallel execution

OPENFOAM® is a registered trademark of OpenCFD Limited.

1.2 Building the Library

```

1 # Navigate to source directory
2 cd openfoam-spaceTimeWindow
3
4 # Build library and utilities
5 # Auto-detects libzstd (compression) and libsodium (encryption)
6 ./Allwmake

```

Listing 1: Building spaceTimeWindow

Optional dependencies are auto-detected by `./Allwmake`:

- **libzstd**: Enables integrated zstd compression of DVZ/DVZT files ($\sim 10\times$ additional reduction)
- **libsodium**: Enables X25519 encryption of boundary data

For manual building with optional features:

```

1 cd src/spaceTimeWindow
2 export FOAM_USE_ZSTD=1 FOAM_USE_SODIUM=1 && wmake
3
4 cd applications/utilities/preProcessing/spaceTimeWindowInitCase
5 export FOAM_USE_ZSTD=1 FOAM_USE_SODIUM=1 && wmake
6
7 cd ../../spaceTimeWindowKeygen
8 export FOAM_USE_SODIUM=1 && wmake

```

Listing 2: Manual build with optional features

2 Concept: Space-Time Window Extraction and Reconstruction

This section introduces the fundamental concepts behind space-time window extraction and reconstruction. Understanding these concepts is essential before diving into the technical implementation details.

2.1 The Challenge of Large-Scale CFD Simulations

Modern computational fluid dynamics (CFD) simulations, particularly Large Eddy Simulations (LES) and Direct Numerical Simulations (DNS), generate enormous amounts of data. A typical industrial LES simulation might involve:

- Meshes with 10^7 – 10^9 cells
- Thousands of timesteps
- Multiple field variables (velocity, pressure, turbulence quantities)
- Terabytes to petabytes of raw output data

Storing and post-processing this data is often impractical. The space-time window approach addresses this by extracting only the data needed to reconstruct the flow in a *region of interest* at a later time.

2.2 The Space-Time Window Concept

A **space-time window** is a bounded region in both space (a 3D box) and time (a time interval) from which we extract the information necessary to reconstruct the flow field. The key insight is that we only need:

1. **Initial conditions**: The complete field state at one timestep (t_0)
2. **Boundary data**: Time-varying values on the boundary of the spatial region

Figure 1 illustrates this concept. The extraction box (blue) defines a spatial subset of the full domain (gray). During simulation, we record the field values on the boundary of this box at every timestep, creating a complete “boundary history” that can later drive a reconstruction simulation.




Figure 1: Space-time window extraction: A bounding box (blue) defines the region of interest within the full computational domain (gray). Boundary data (orange) is recorded at every timestep as the flow passes through.

2.3 Extraction and Reconstruction Workflow

The workflow consists of two distinct phases that can be separated in time and computational resources, as shown in fig. 2.




Figure 2: Two-phase workflow: The extraction phase runs on HPC resources and produces compact data files. The reconstruction phase can run on modest hardware, reproducing the flow in the region of interest.

2.3.1 Extraction Phase (HPC)

During the original simulation on high-performance computing resources:

1. The `spaceTimeWindowExtract` function object monitors the simulation
2. At the start time (t_0), it extracts the complete field state within the bounding box
3. At every timestep, it interpolates field values to the boundary faces and stores them
4. The output is a compact dataset: initial conditions + time-varying boundary data

2.3.2 Reconstruction Phase (Offline)

Later, on more modest computational resources:

1. The `spaceTimeWindowInitCase` utility creates a complete simulation case
2. The subset mesh is extracted from the original
3. Boundary conditions are configured to read the stored boundary data
4. The solver runs, reproducing the flow within the extraction box

 Tip

The reconstruction phase typically requires 100–10,000× fewer computational resources than the original simulation, making it feasible to run on a workstation or laptop.

2.4 Parallel Execution with MPI Domain Decomposition

Large-scale CFD simulations invariably run in parallel using MPI (Message Passing Interface). The computational domain is divided among multiple processors, each handling a subset of the mesh. The `spaceTimeWindow` library handles both serial and parallel execution transparently.

Figure 3 illustrates how domain decomposition affects the extraction process. The extraction box may span multiple processor domains, requiring coordination to gather boundary data from all relevant processors.

2.4.1 How Parallel Extraction Works

In parallel mode, the extraction process operates as follows:

1. **Local identification:** Each processor identifies which of its cells fall within the extraction bounding box
2. **Local extraction:** Each processor extracts field values for its local portion of the boundary
3. **Global gathering:** Boundary data is gathered to the master processor using MPI collective operations
4. **Unified output:** The master processor writes a single, consolidated boundary data file for each timestep

Figure 4 shows the detailed data flow during the extraction process. The `spaceTimeWindowExtract` function object runs on each MPI process, extracting local contributions that are then gathered and merged into unified output files.

This approach ensures that:

Parallel Extraction: Domain Decomposition with 8 MPI Processes




Figure 3: Domain decomposition in parallel simulations. The extraction box (blue) spans multiple processor domains. Each processor extracts its local contribution, and data is gathered to the master process for writing unified boundary data files.

MPI Data Gathering by spaceTimeWindowExtract




Figure 4: Detailed MPI data gathering during parallel extraction. Each processor runs the `spaceTimeWindowExtract` function object, which identifies local boundary faces, interpolates field values to face centers, and participates in a collective gather operation. The master processor (P0) writes the merged data to a single file per timestep.


- The reconstruction case is **decomposition-independent**—the same boundary data works regardless of how the original simulation was parallelized
- Output files are compact and self-contained, with no processor-specific information
- The reconstruction can run in serial or with a different parallel decomposition

i Note

In parallel mode, mesh information is not written during extraction. Instead, the extraction box parameters are stored, and `spaceTimeWindowInitCase` creates the subset mesh from reconstructed fields. Run `reconstructPar -time <startTime>` on the source case before initializing the reconstruction case.

2.5 Data Storage: Initial Fields and Boundary Data

The extracted data consists of two distinct components, illustrated in fig. 5.



→ Boundary data dominates storage. Use dvzt format with integrated zstd for 300×+ compression.

Figure 5: Data storage structure: Initial fields are written once at t_0 , while boundary data is accumulated at every timestep. Compression (DVZ/DVZT + integrated zstd) is critical for the boundary data component.

2.5.1 Initial Fields

The initial fields provide the starting state for reconstruction:

- **What:** Complete volumetric data (U, p , turbulence quantities) within the extraction box
- **When:** Written once at the extraction start time (t_0)
- **Format:** Standard OpenFOAM field format (ASCII)
- **Size:** Proportional to number of cells in the subset, typically megabytes

2.5.2 Boundary Data

The boundary data provides the time-varying boundary conditions:

- **What:** Face-interpolated field values on the boundary of the extraction box
- **When:** Written at every timestep throughout the simulation
- **Format:** Configurable (ASCII, binary, or compressed dvz/dvzt)
- **Size:** Proportional to (number of boundary faces) \times (number of timesteps)

 **Warning**

Boundary data accumulates rapidly during long simulations. For a typical LES with 10^4 boundary faces and 10^4 timesteps, uncompressed ASCII data can reach tens of gigabytes. The dvzt compression format reduces this by 30–50 \times , and with integrated zstd compression the reduction reaches 300 \times +

2.6 Why Reconstruction Works

The reconstruction is possible because the governing equations (Navier-Stokes for incompressible flow) are well-posed with appropriate boundary conditions. Mathematically, if we denote:

- $\Omega \subset \mathbb{R}^3$: the extraction box domain
- $\partial\Omega$: the boundary of the extraction box
- $\mathbf{u}(\mathbf{x}, t)$: the velocity field
- $p(\mathbf{x}, t)$: the pressure field

Then knowing $\mathbf{u}(\mathbf{x}, t_0)$ for $\mathbf{x} \in \Omega$ (initial condition) and $\mathbf{u}(\mathbf{x}, t)$ for $\mathbf{x} \in \partial\Omega, t \in [t_0, t_n]$ (boundary data) is sufficient to uniquely determine $\mathbf{u}(\mathbf{x}, t)$ for all $\mathbf{x} \in \Omega, t \in [t_0, t_n]$.

The reconstruction achieves “bit-reproducible” results when:

1. Time interpolation is disabled (`timeInterpolationScheme none`)
2. The reconstruction runs with the same timestep as extraction
3. The same solver settings and numerical schemes are used

With time interpolation enabled (recommended for practical use), the reconstruction closely approximates the original solution, with errors controlled by the interpolation scheme (linear or cubic Catmull-Rom spline).

2.7 Automatic Boundary Condition Configuration

A key feature of the spaceTimeWindow library is the **automatic configuration** of boundary conditions during offline reconstruction. When you run `spaceTimeWindowInitCase`, the utility analyzes the extracted data and configures appropriate boundary conditions for each field—no manual setup required.

Figure 6 illustrates the automatic configuration process that occurs when setting up a reconstruction case on a laptop or desktop computer.




Figure 6: Recommended boundary condition configuration for pressure-coupled reconstruction. The `spaceTimeWindowInitCase` utility can assign BCs based on available data: `spaceTimeWindow` for velocity (Dirichlet), `spaceTimeWindowCoupledPressure` for pressure (mixed Dirichlet/Neumann), and `zeroGradient` for turbulence fields.

2.7.1 Resolution-Independent Reconstruction

The reconstruction can run at **different mesh resolutions** than the original extraction. This enables several important use cases:

- **Coarser mesh:** Faster turnaround for exploratory analysis or debugging
- **Finer mesh:** Higher resolution studies without re-running the expensive global simulation
- **Same resolution:** Bit-reproducible reconstruction for validation

Figure 7 shows how the `spaceTimeWindow` boundary condition automatically handles resolution differences at runtime.

2.7.2 Time Interpolation at Runtime

Similarly, the boundary condition handles **time interpolation** automatically when the reconstruction timestep differs from the extraction timestep (which is common with adaptive timestepping). Figure 8 illustrates the temporal interpolation process.

✓ Tip

The combination of automatic spatial and temporal interpolation means you can:

- Run reconstruction on a laptop with a coarser mesh for quick visualization
- Use adaptive timestepping (`adjustTimeStep yes`) without worrying about timestep alignment




Figure 7: Resolution-independent reconstruction: The spaceTimeWindow boundary condition automatically handles mesh resolution differences between extraction and reconstruction. Spatial interpolation is performed on-the-fly during the simulation.




Figure 8: Time interpolation during reconstruction. When the solver requests field values at times not exactly matching stored timesteps, the boundary condition interpolates using linear (2 points) or cubic Catmull-Rom (4 points) schemes. Cubic interpolation provides smooth derivatives, essential for adaptive timestepping.

- Later refine the mesh for detailed analysis without re-extracting data

All interpolation happens transparently at runtime—no preprocessing steps required.

3 Encryption of Boundary Data

The spaceTimeWindow library supports X25519 asymmetric encryption using libsodium sealed boxes. This allows secure distribution of simulation data where:

- Extraction uses only the **public key**
- Reconstruction requires the **private key**
- Data cannot be decrypted without the private key

3.1 Key Generation

Generate a keypair using the provided utility:

```

1 spaceTimeWindowKeygen
2 # Output:
3 # Public key: fqzYQ0U8j27tFEr5WzEMylbvXYP+9CAyk0JhwwZ2rwg=
4 # Private key: ****H8exb*****n4Kv0nu5gqljI2RPBCwl4=
5 #
6 # IMPORTANT: Store the private key securely!

```

Listing 3: Generating encryption keys

⚠ Warning

The private key must be stored securely and never committed to version control. Anyone with the private key can decrypt all boundary data encrypted with the corresponding public key.

3.2 Secure Key Storage

The private key should be stored using one of the following methods:

- **Hardware security keys:** YubiKey, Nitrokey, or similar FIDO2/PIV-capable devices provide the strongest protection. The private key never leaves the hardware token.
- **Secure password managers:** KeePassXC, 1Password, Bitwarden, or similar tools with strong master passwords and optional hardware key integration.
- **Encrypted vaults:** GPG-encrypted files or OS keychains (macOS Keychain, GNOME Keyring, Windows Credential Manager).

✓ Tip

For maximum security, generate the keypair on an air-gapped machine, store the private key on a hardware token, and only transfer the public key to the HPC system.

3.3 Encrypted Extraction

Add the public key to your extraction configuration in `system/controlDict`:

```

1  functions
2  {
3      extractSubset
4      {
5          type          spaceTimeWindowExtract;
6          libs          (spaceTimeWindow);
7
8          box           ((0.05 -0.25 0.01) (0.90 0.25 0.38));
9          outputDir    "../subset-case";
10         fields        (U p nut);
11
12         writeFormat   deltaVarint;
13
14         // Enable encryption with public key
15         publicKey     "fqzYQ0U8j27tFEr5WzEMylbvXYP+9CAyk0JhwwZ2rwg=";
16
17         writeControl   timeStep;
18         writeInterval 1;
19     }
20 }
```

Listing 4: Encrypted extraction configuration

Encrypted files have the `.enc` extension appended to the codec extension (e.g., `U.dvz.enc`, or `U.dvz.zstd.enc` when integrated zstd compression is also enabled).

3.4 Decryption During Case Initialization

When running `spaceTimeWindowInitCase` on encrypted data, you will be prompted for the private key:

```

1 cd subset-case
2 spaceTimeWindowInitCase -sourceCase ../source-case
3 # Enter private key (base64): [input not echoed]
```

Listing 5: Decrypting boundary data

The utility automatically:

1. Derives the public key from the private key
2. Decrypts all `.enc` files in `boundaryData`
3. Removes encrypted files after successful decryption

3.5 Security Properties

- **Sealed box encryption:** Anonymous sender, only recipient can decrypt
- **X25519 key exchange:** 128-bit security level

- **Xsalsa20-Poly1305:** Authenticated encryption
- **No plaintext on disk:** Encryption occurs before writing

3.5.1 Cryptographic Foundations

The encryption uses elliptic curve Diffie-Hellman on Curve25519. Key generation produces a private scalar s and public point (eq. (1)):

$$P_{\text{pub}} = s \cdot G \quad (1)$$

where G is the curve's base point. The discrete logarithm problem makes recovering s from P_{pub} computationally infeasible.

Sealed Box Construction For each encryption, a sealed box generates an ephemeral keypair $(e, e \cdot G)$ and computes a shared secret:

$$K = \text{BLAKE2b}(e \cdot P_{\text{recv}}) = \text{BLAKE2b}(e \cdot s_{\text{recv}} \cdot G) \quad (2)$$

The ciphertext is constructed as shown in eq. (3):

$$C = (e \cdot G) \| \text{Xsalsa20-Poly1305}_K(M) \quad (3)$$

where M is the plaintext boundary data and $\|$ denotes concatenation. The ephemeral public key $(e \cdot G)$ is prepended to allow decryption.

Decryption The recipient computes the same shared secret using their private key (eq. (4)):

$$K = \text{BLAKE2b}(s_{\text{recv}} \cdot (e \cdot G)) \quad (4)$$

By the associativity of scalar multiplication on elliptic curves (eq. (5)):

$$s_{\text{recv}} \cdot (e \cdot G) = e \cdot (s_{\text{recv}} \cdot G) = e \cdot P_{\text{recv}} \quad (5)$$

This ensures both parties derive the same key K without ever transmitting it.

Security Level X25519 provides approximately 128 bits of security. Breaking the encryption requires either:

- Solving the elliptic curve discrete logarithm problem (ECDLP): $O(2^{128})$ operations
- Brute-forcing the symmetric key: $O(2^{256})$ operations for Xsalsa20

3.6 Ethical Considerations and HPC Transparency

Most high-performance computing (HPC) centers require full transparency regarding the simulations performed on their infrastructure. For this reason, **encryption support is an optional compile-time feature** controlled by the `FOAM_USE_SODIUM` flag. Centers may choose not to enable this functionality.

It is important to understand what encryption does and does not protect:

- **What is protected:** The time-varying boundary field data in `constant/boundaryData`. When downloaded, this data is sealed and cannot be read without the private key.

- **What is NOT protected:** The test case setup (mesh, dictionaries, initial conditions) remains unencrypted. Anyone can re-run the global simulation to regenerate the boundary data—encryption does not prevent this.

The practical security model relies on the following observations:

1. Re-running the global simulation to obtain internal fields is **computationally expensive** and requires HPC resources. Such jobs are recorded in the scheduler logs, providing an audit trail.
2. Offline recalculation on a personal workstation is typically infeasible due to computational cost.
3. If file permissions on HPC scratch directories are misconfigured and the user chooses not to write timestep data from the global simulation, other users can only copy the unencrypted test case setup—not the valuable boundary data.

i Note

The encryption feature is designed for protecting intellectual property during data distribution, not for hiding simulations from HPC administrators. Always comply with your computing center's acceptable use policies.

3.7 Protection Against Malware and Ransomware

The spaceTimeWindow approach offers an additional security benefit: by storing only the minimal data required for reconstruction, the valuable simulation results can be protected against malware and ransomware attacks.

3.7.1 Secure Archival Strategy

The reconstruction case contains only:

- The subset mesh (`constant/polyMesh`)
- Encrypted boundary data (`constant/boundaryData/*.zstd.enc` or `*.enc`)
- Initial fields for one timestep
- Configuration dictionaries

This minimal dataset (typically a few gigabytes) can be stored on:

- **Write-Once Read-Many (WORM) media:** Optical discs (M-DISC), tape archives with WORM capability, or cloud storage with object lock/immutability settings (AWS S3 Object Lock, Azure Immutable Blob Storage).
- **Secure vaults:** Hardware-encrypted drives kept offline, or institutional secure storage with access controls and audit logging.
- **Air-gapped backups:** Disconnected storage that cannot be reached by network-based malware.

3.7.2 Security Model

1. **Encrypted boundary data is immutable:** Once written to WORM storage, the encrypted files cannot be modified or deleted by malware.
2. **Private key stored separately:** The decryption key resides on a hardware token or secure password manager, not on the same system as the data.
3. **Reconstruction is always possible:** Even if the HPC system or workstation is compromised, the archived data remains intact and can be decrypted on a clean system.
4. **Minimal attack surface:** The small archive size makes backup and verification practical, unlike terabyte-scale full simulation outputs.

 Tip

For critical simulations, archive the encrypted reconstruction case to WORM storage immediately after extraction completes. The full simulation data can then be deleted from HPC scratch space, eliminating the risk of data loss from ransomware or accidental deletion.

4 Boundary Condition Approaches

The library provides three approaches for applying boundary conditions on the extraction boundary (`oldInternalFaces`):

4.1 Pressure-Coupled BC – Recommended

The **recommended approach** for accurate reconstruction uses the `spaceTimeWindowCoupledPressure` boundary condition for pressure, which properly handles the elliptic nature of the pressure equation.

- **Velocity (U):** `spaceTimeWindow` — pure Dirichlet from extracted data
- **Pressure (p):** `spaceTimeWindowCoupledPressure` — mixed Dirichlet/Neumann based on flux
- **Turbulence fields:** `zeroGradient`

Why this approach is physically correct:

- Pressure is governed by an elliptic Poisson equation — cutting the domain severs the pressure's ability to “feel” the flow outside the window
- The BC uses smooth tanh-based blending to transition between Dirichlet and Neumann limits
- Two blending modes are available:
 - `fluxMagnitude` (default, recommended): Dirichlet at stagnant faces, Neumann at active flow faces — more compatible with continuity
 - `flowDirection`: Dirichlet at inflow faces, Neumann at outflow faces — the original approach

- This **reproduces the original flow** rather than computing a different solution

```

1 // In 0/U
2 oldInternalFaces
3 {
4     type          spaceTimeWindow;
5     dataDir       "constant/boundaryData";
6     allowTimeInterpolation  true;
7     timeInterpolationScheme cubic;
8     value         uniform (0 0 0);
9 }
10
11 // In 0/p
12 oldInternalFaces
13 {
14     type          spaceTimeWindowCoupledPressure;
15     dataDir       "constant/boundaryData";
16     phi           phi;
17     blendingMode fluxMagnitude; // or flowDirection
18     dirichletWeight 0.8;        // Dirichlet strength (0-1)
19     neumannWeight  0.8;        // Neumann strength (0-1)
20     transitionWidth 0.1;      // Blending sharpness
21     allowTimeInterpolation  true;
22     timeInterpolationScheme cubic;
23     value         uniform 0;
24 }
```

Listing 6: Pressure-coupled boundary conditions

i Note

The pressure-coupled approach requires extracting pressure gradients during extraction. Enable this with `extractGradients true` and `gradientFields (p)` in the extraction configuration. See section 6.1 for details.

4.2 Inlet-Outlet BC (-inletOutletBC) – For Quick Tests

Uses `spaceTimeWindowInletOutlet` BC which applies:

- **Velocity (U):** Flux-based switching – Dirichlet (prescribed value) at inflow faces, zero-Gradient at outflow faces
- **All scalar fields (p, nut, k, epsilon, omega):** zeroGradient

This approach is useful for quick tests and situations where exact reproduction of the original flow is not required:

- Turbulent flows have instantaneous velocity fluctuations that can reverse direction locally
- The flux-based switching automatically adapts to the instantaneous flow direction at each face

- However, pressure uses zeroGradient, which can lead to solution drift over time

```
1 spaceTimeWindowInitCase -sourceCase ../source-case -inletOutletBC
```

Listing 7: Using inlet-outlet BC

△ Warning

The inlet-outlet BC does not reproduce the original pressure field. The subset simulation will evolve to a **different solution** because pressure information from outside the extraction window is lost. For accurate reconstruction, use the pressure-coupled approach instead.

4.3 Fixed Outlet Direction (-outletDirection)

Creates a separate outlet patch from faces at one edge of the extraction box:

- **oldInternalFaces**: Pure Dirichlet BC (spaceTimeWindow) on all remaining faces
- **outlet**: Pressure relief patch with inletOutlet for U and zeroGradient for p

Use when the mean flow direction is well-defined and outflow always occurs at a known location.

```
1 spaceTimeWindowInitCase -sourceCase ../source-case -outletDirection "(1 0 0)"
```

Listing 8: Using fixed outlet direction

△ Warning

These options (-inletOutletBC and -outletDirection) are mutually exclusive. Using both together produces an error with guidance.

4.4 Boundary Condition Assignment Logic

For each field in the initial time directory, `spaceTimeWindowInitCase` assigns BCs according to Table 1.

Table 1: Boundary condition assignment logic

In boundaryData?	-inletOutletBC?	Field Type	BC Applied
Yes	Yes	vector (U)	spaceTimeWindowInletOutlet
Yes	Yes	scalar	zeroGradient
Yes	No	any	spaceTimeWindow (Dirichlet)
No	any	any	zeroGradient

This ensures fields without time-varying data automatically get appropriate BCs.

5 Workflow Overview

The space-time window reconstruction workflow is strictly sequential, as shown in fig. 9.




Figure 9: Space-time window reconstruction workflow

5.1 Serial Workflow

```

1 # 1. Run extraction during simulation
2 cd source-case
3 pimpleFoam # With spaceTimeWindowExtract function object
4
5 # 2. Initialize reconstruction case (recommended: inlet-outlet BC)
6 cd ../subset-case
7 spaceTimeWindowInitCase -sourceCase ../source-case -inletOutletBC
8
9 # 3. Run reconstruction
10 pimpleFoam
  
```

Listing 9: Complete serial workflow

In serial mode:

- Mesh and initial fields are written directly to the output directory
- Boundary data is written at every timestep
- The extractor forces field writes at t_1 and t_2 (for interpolation buffers)

i Note

In serial mode, the extractor writes internal fields only at t_0 (extraction start). The t_1 and t_2 internal fields are **not** automatically saved. If you need linear interpolation starting at t_1 , ensure your `writeInterval` captures t_1 , or use parallel mode which forces writes at t_0 , t_1 , and t_2 .

5.2 Parallel Workflow

1. Run the original simulation in parallel with extraction enabled
2. Run `reconstructPar` for the desired start timestep (t_2 for cubic interpolation)
3. Run `spaceTimeWindowInitCase -sourceCase <path> -inletOutletBC`
4. Run the solver on the subset case (serial or parallel)

```

1 # Step 1: Run parallel extraction
2 cd source-case
3 mpirun -np 4 pimpleFoam -parallel
4 # Extraction automatically forces writes at t_0, t_1, t_2
5
6 # Step 2: Reconstruct fields at cubic-safe start time (t_2)
7 reconstructPar -time 0.0002
8
9 # Step 3: Initialize subset case (recommended: inlet-outlet BC)
10 cd ../subset-case
11 spaceTimeWindowInitCase -sourceCase ../source-case -inletOutletBC
12
13 # Step 4: Run reconstruction
14 pimpleFoam

```

Listing 10: Complete parallel workflow

In parallel mode:

- Boundary data from all processors is gathered to master and written as single files
- The extractor forces field writes at t_0 , t_1 , and t_2 for interpolation buffers (t_0 only in parallel)
- Only extraction box parameters are written (not mesh) — `spaceTimeWindowInitCase` creates the mesh
- `reconstructPar` must be run at the reconstruction start time (t_2) to provide source fields

6 Configuration Reference

6.1 Extraction Configuration

The extraction function object is configured in `system/controlDict` within the `functions` sub-dictionary. The available parameters are listed in table 2.

```

1 functions
2 {
3     extractSubset
4     {
5         type          spaceTimeWindowExtract;
6         libs          (spaceTimeWindow);
7
8         // Bounding box: ((minX minY minZ) (maxX maxY maxZ))
9         // Must be fully internal to domain
10        box           ((0.05 -0.25 0.01) (0.90 0.25 0.38));
11
12        outputDir      "../subset-case";
13
14        // Fields for time-varying boundary data (written every timestep)
15        fields         (U p);           // U and p for pressure-coupled BC
16

```

```

17   // Fields for initial conditions (optional, defaults to 'fields')
18   initialFields  (U p nut);           // nut needed for turbulence model
19
20   // Gradient extraction for pressure-coupled BC (RECOMMENDED)
21   extractGradients true;
22   gradientFields  (p);             // Extract normal gradient of p
23
24   // Write format: ascii, binary, deltaVarint, or dvzt
25   writeFormat      dvzt;          // Recommended: best compression
26
27   // Precision for deltaVarint/dvzt (default: 6)
28   deltaVarintPrecision 6;
29
30   // Keyframe interval for dvzt (default: 20)
31   dvztKeyframeInterval 20;
32
33   // Gzip compression (ignored for deltaVarint/dvzt)
34   writeCompression off;
35
36   // Optional encryption
37   // publicKey      "base64-encoded-public-key";
38
39   // Time window (optional)
40   // timeStart     0.0;
41   // timeEnd       1.0;
42
43   writeControl    timeStep;
44   writeInterval   1;
45 }
46 }
```

Listing 11: Complete extraction configuration

6.2 Initial Fields vs Boundary Data Fields

The extraction can separate which fields are used for:

- **Initial conditions (initialFields):** Fields written to the start time directory for solver initialization
- **Time-varying boundary data (fields):** Fields written to boundaryData/ for time-varying BCs

For the recommended pressure-coupled approach:

```

1 // In extraction function object
2 fields        (U p);           // U and p for time-varying BC
3 initialFields (U p nut);       // nut needed for turbulence model
4 extractGradients true;         // Enable gradient extraction
5 gradientFields (p);           // Extract normal gradient of p (creates
                                gradp)
```

Listing 12: Field selection strategy for pressure-coupled BC (RECOMMENDED)

Table 2: Extraction parameters

Parameter	Type	Required	Description
box	pointPair	Yes	Extraction bounding box
outputDir	fileName	Yes	Output case directory
fields	wordList	Yes	Fields for time-varying BC (boundaryData)
initialFields	wordList	No	Fields for initial conditions (default: same as fields)
extractGradients	bool	No	Enable gradient extraction (default: false)
gradientFields	wordList	No	Fields for gradient extraction (requires extractGradients)
writeFormat	word	No	ascii, binary, deltaVarint, or dvzt
deltaVarintPrecision	label	No	Decimal digits (default: 6)
dvztKeyframeInterval	label	No	Keyframe interval for dvzt (default: 20)
writeCompression	switch	No	Gzip compression (ignored for deltaVarint/dvzt)
publicKey	string	No	Base64 public key for encryption
timeStart	scalar	No	Extraction start time
timeEnd	scalar	No	Extraction end time

Note: nut is only needed in `initialFields` (not in `fields`) because turbulent viscosity is computed by the turbulence model at runtime and uses zeroGradient BC on `oldInternalFaces`.

Fields NOT in `boundaryData` automatically get zeroGradient BC on `oldInternalFaces`.

✓ Tip

For accurate reconstruction, use the **pressure-coupled BC approach**: extract both U and p with gradient extraction enabled for p. This allows the `spaceTimeWindowCoupledPressure` BC to properly constrain the pressure field and reproduce the original solution.

i Note

When `extractGradients` is enabled, the extractor computes the face-normal gradient at each boundary face using the equation:

$$\left. \frac{\partial \phi}{\partial n} \right|_f = \frac{\phi_N - \phi_P}{|\mathbf{d}_{PN}|}$$

where ϕ_P is the value at the face owner cell, ϕ_N is the value at the neighbour cell, and \mathbf{d}_{PN} is the vector from owner to neighbour cell centre. The gradient data is written to files named `grad<fieldName>` (e.g., `gradp`) in the `boundaryData` directory.

6.3 Boundary Condition Configuration

6.3.1 spaceTimeWindowCoupledPressure (Recommended for Pressure)

Mixed (Robin) boundary condition for pressure that blends Dirichlet and Neumann conditions based on local flux. This is the **recommended approach** for pressure at extraction boundaries. Two blending modes are available:

- **fluxMagnitude** (default, recommended): Dirichlet at stagnant faces (zero flux), Neumann at active flow faces — better continuity compatibility
- **flowDirection**: Dirichlet at inflow, Neumann at outflow — the original flow-direction-based approach

The parameters are listed in table 3.

```

1  boundaryField
2  {
3      oldInternalFaces
4      {
5          type                  spaceTimeWindowCoupledPressure;
6          dataDir               "constant/boundaryData";
7          phi                   phi;           // Flux field name
8          blendingMode         fluxMagnitude; // or flowDirection
9          dirichletWeight      0.8;          // Dirichlet strength (0-1)
10         neumannWeight       0.8;          // Neumann strength (0-1)
11         transitionWidth     0.1;          // Blending sharpness
12         allowTimeInterpolation true;
13         timeInterpolationScheme cubic;    // or "linear" or "none"
14         value                 uniform 0;
15     }
16 }
```

Listing 13: spaceTimeWindowCoupledPressure configuration

How it works:

1. Reads pressure values and pressure gradients from boundaryData
2. Computes local flux at each face
3. Calculates blending weight using smooth tanh transition based on the selected mode:
 - **fluxMagnitude**: $w = w_D + (w'_N - w_D) \cdot \tanh(|\phi_f|/\phi_{max}/\sigma)$ — where w_D is `dirichletWeight` and $w'_N = 1 - w_N$
 - **flowDirection**: $w = \frac{1}{2} ((w_D + w'_N) - (w_D - w'_N) \cdot \tanh(\phi_f/\phi_{max}/\sigma))$
4. In `fluxMagnitude` mode: at **stagnant faces** ($w \rightarrow w_D$) applies Dirichlet; at **high-flux faces** ($w \rightarrow w'_N$) applies Neumann
5. In `flowDirection` mode: at **inflow faces** ($w \rightarrow w_D$) applies Dirichlet; at **outflow faces** ($w \rightarrow w'_N$) applies Neumann

The value coefficients are:

$$\text{valueFraction} = w \cdot w_D + (1 - w) \cdot 0 = w \cdot w_D \quad (6)$$

$$\text{refValue} = p_{\text{extracted}} \quad (7)$$

$$\text{refGrad} = \left. \frac{\partial p}{\partial n} \right|_{\text{extracted}} \quad (8)$$

where w_D is `dirichletWeight` and the gradient coefficient is $(1 - w) \cdot w_N$ with w_N being `neumannWeight`.

Table 3: spaceTimeWindowCoupledPressure parameters

Parameter	Type	Default	Description
<code>dataDir</code>	<code>fileName</code>	<code>constant/boundaryData</code>	Path to boundary data
<code>phi</code>	<code>word</code>	<code>phi</code>	Name of flux field
<code>blendingMode</code>	<code>word</code>	<code>fluxMagnitude</code>	<code>fluxMagnitude</code> or <code>flowDirection</code>
<code>dirichletWeight</code>	<code>scalar</code>	0.8	Dirichlet contribution strength (0-1)
<code>neumannWeight</code>	<code>scalar</code>	0.8	Neumann contribution strength (0-1)
<code>transitionWidth</code>	<code>scalar</code>	0.1	Sharpness of flux transition
<code>allowTimeInterpolation</code>	<code>bool</code>	<code>false</code>	Permit interpolation for missing timesteps
<code>timeInterpolationScheme</code>	<code>word</code>	<code>linear</code>	<code>none</code> , <code>linear</code> , or <code>cubic</code>

i Note

This BC requires gradient data in `boundaryData`. Enable gradient extraction with `extractGradients true` and `gradientFields (p)` in the extraction configuration. The gradient file must be named `gradp` in each timestep directory.

6.3.2 spaceTimeWindowInletOutlet

Flux-based boundary condition that reads pre-computed velocity values and applies them only at inflow faces. The parameters are listed in table 4.

```

1 boundaryField
2 {
3     oldInternalFaces
4     {
5         type          spaceTimeWindowInletOutlet;
6         dataDir       "constant/boundaryData";
7         phi           phi;           // Flux field name
8         allowTimeInterpolation  true;
9         timeInterpolationScheme cubic;        // or "linear" or "none"
10        value         uniform (0 0 0);

```

```

11     }
12 }
```

Listing 14: spaceTimeWindowInletOutlet configuration**How it works:**

1. At each timestep, reads velocity values from boundaryData (with optional time interpolation)
2. Computes flux through each face: $\phi_f = \mathbf{U}_f \cdot \mathbf{S}_f$
3. For inflow faces ($\phi < 0$): applies prescribed velocity from boundaryData
4. For outflow faces ($\phi \geq 0$): applies zeroGradient (extrapolates from interior)

Table 4: spaceTimeWindowInletOutlet parameters

Parameter	Type	Default	Description
dataDir	fileName	constant/boundaryData	Path to boundary data
phi	word	phi	Name of flux field
allowTimeInterpolation	bool	false	Permit interpolation for missing timesteps
timeInterpolationScheme	word	linear	none, linear, or cubic

6.3.3 spaceTimeWindow (Pure Dirichlet)

Pure Dirichlet boundary condition that prescribes values on all faces regardless of flow direction. The parameters are listed in table 5.

```

1 boundaryField
2 {
3     oldInternalFaces
4     {
5         type          spaceTimeWindow;
6         dataDir       "constant/boundaryData";
7         fixesValue    true;           // Tells adjustPhi these values
8                     are fixed
9         allowTimeInterpolation  true;
10        timeInterpolationScheme cubic;
11        reportFlux      true;
12        value          uniform (0 0 0);
13    }
14 }
```

Listing 15: spaceTimeWindow configuration

Use this for:

- Scalar fields when not using `-inletOutletBC`
- Situations where fixed values are explicitly desired on all faces

Table 5: spaceTimeWindow parameters

Parameter	Type	Default	Description
dataDir	fileName	constant/boundaryData	Path to boundary data
fixesValue	bool	true	Report to <code>adjustPhi</code> that values are fixed
allowTimeInterpolation	bool	false	Enable time interpolation
timeInterpolationScheme	word	linear	none, linear, or cubic
reportFlux	bool	false	Print net flux through patch (velocity only)
setAverage	bool	false	Adjust field average
offset	Type	Zero	Offset value

6.4 Case Initialization Options

The `spaceTimeWindowInitCase` utility accepts command-line options (see table 6):

```

1 # Recommended: inlet-outlet BC for unsteady turbulent flows
2 spaceTimeWindowInitCase -sourceCase ../source-case -inletOutletBC
3
4 # Alternative: fixed outlet direction for steady-mean flows
5 spaceTimeWindowInitCase -sourceCase ../source-case -outletDirection "(1 0 0)"
6
7 # With mass flux correction (optional, ensures exact mass conservation)
8 spaceTimeWindowInitCase -sourceCase ../source-case -inletOutletBC
    -correctMassFlux

```

Listing 16: spaceTimeWindowInitCase usage**Table 6:** spaceTimeWindowInitCase options

Option	Type	Description
<code>-sourceCase</code>	directory	Source case where extraction ran (required)
<code>-extractDir</code>	directory	Directory with extracted data (default: cwd)
<code>-inletOutletBC</code>	flag	Use flux-based inlet-outlet BC for U, zeroGradient for scalars (for quick tests)
<code>-outletDirection</code>	vector	Create fixed outlet patch in given direction (e.g., "(1 0 0)")
<code>-outletFraction</code>	scalar	Fraction of box extent for outlet region (default: 0.1)
<code>-correctMassFlux</code>	flag	Apply least-squares mass flux correction to boundaryData
<code>-initialFields</code>	list	Override initial fields list (e.g., "(U p nut k)")
<code>-refineLevel</code>	label	Refine mesh N times (spatial interpolation at runtime)
<code>-coarsenLevel</code>	label	Coarsen mesh N times (spatial interpolation at runtime)
<code>-overwrite</code>	flag	Overwrite existing files

△ Warning

-inletOutletBC and -outletDirection are mutually exclusive. Using both together produces an error with guidance on which option to choose.

6.5 Mesh Coarsening and Refinement

The reconstruction mesh can be coarsened or refined relative to the extraction mesh (table 7). This enables running reconstructions at different resolutions than the original simulation.

```

1 # Refine mesh (finer than extraction)
2 spaceTimeWindowInitCase -sourceCase ../source -inletOutletBC -refineLevel 1
3
4 # Coarsen mesh (coarser than extraction)
5 spaceTimeWindowInitCase -sourceCase ../source -inletOutletBC -coarsenLevel 1
6
7 # Multiple levels
8 spaceTimeWindowInitCase -sourceCase ../source -inletOutletBC -refineLevel 2

```

Listing 17: Mesh resolution options

Table 7: Mesh resolution options

Option	Type	Description
-refineLevel	label	Refine mesh N times (each level splits cells $\times 8$)
-coarsenLevel	label	Coarsen mesh N times (each level merges cells)

△ Warning

-refineLevel and -coarsenLevel are mutually exclusive. Use only one at a time.

6.5.1 Spatial Interpolation Algorithms

When the reconstruction mesh differs from the extraction mesh, spatial interpolation is required for boundary data. The spaceTimeWindow boundary conditions handle this automatically at runtime using different algorithms depending on the resolution change.

Refinement: Barycentric Interpolation with 2D Delaunay Triangulation When the target mesh has more faces than the source (refinement), the algorithm triangulates source face centers using the Bowyer-Watson algorithm, as illustrated in fig. 10. For each target face center, the enclosing triangle is found and barycentric weights are computed. If the target point lies outside all triangles (which can happen because the extracted submesh uses original cells from the source case and may have irregular boundaries), the algorithm finds the nearest triangle by centroid distance and uses clamped barycentric coordinates. This provides smooth C^0 continuous interpolation.

Given a target point p inside a triangle with vertices v_1, v_2, v_3 , the barycentric coordinates $(\lambda_1, \lambda_2, \lambda_3)$ satisfy eq. (9):

$$p = \lambda_1 v_1 + \lambda_2 v_2 + \lambda_3 v_3, \quad \text{where} \quad \lambda_1 + \lambda_2 + \lambda_3 = 1 \quad (9)$$

The weights are computed from signed triangle areas (eq. (10)):

$$\lambda_1 = \frac{A(\mathbf{p}, \mathbf{v}_2, \mathbf{v}_3)}{A(\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3)}, \quad \lambda_2 = \frac{A(\mathbf{v}_1, \mathbf{p}, \mathbf{v}_3)}{A(\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3)}, \quad \lambda_3 = \frac{A(\mathbf{v}_1, \mathbf{v}_2, \mathbf{p})}{A(\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3)} \quad (10)$$

where the signed area of a 2D triangle is given by eq. (11):

$$A(\mathbf{a}, \mathbf{b}, \mathbf{c}) = \frac{1}{2} [(b_x - a_x)(c_y - a_y) - (c_x - a_x)(b_y - a_y)] \quad (11)$$

The interpolated field value at the target point is then (eq. (12)):

$$\phi(\mathbf{p}) = \lambda_1 \phi_1 + \lambda_2 \phi_2 + \lambda_3 \phi_3 \quad (12)$$

Barycentric Interpolation (Refinement)




Figure 10: Barycentric interpolation for mesh refinement: source face centers (blue) are triangulated, and each target point (green) is interpolated using barycentric weights from the enclosing triangle

Coarsening: Area-Weighted Averaging When the target mesh has fewer faces than the source (coarsening), an octree is built from source points for efficient spatial lookup, as shown in fig. 11. For each target face, all source points within a search radius are found and averaged with equal weights. If no points are found, the search radius is progressively expanded. This ensures conservation of integral quantities.

For a target point \mathbf{p} with search radius r , define the set of contributing source points (eq. (13)):

$$\mathcal{S}(\mathbf{p}, r) = \{i : \|x_i - \mathbf{p}\| \leq r\} \quad (13)$$

The interpolated value is the arithmetic mean of all contributing sources (eq. (14)):

$$\phi(\mathbf{p}) = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \phi_i \quad (14)$$

If $|\mathcal{S}| = 0$, the search radius is expanded by a factor $\alpha > 1$ (typically $\alpha = 1.5$) and the search is repeated (eq. (15)):

$$r \leftarrow \alpha \cdot r \quad \text{until} \quad |\mathcal{S}(\mathbf{p}, r)| > 0 \quad (15)$$

The octree provides $O(\log N)$ lookup complexity for finding points within the search radius, making the algorithm efficient even for large meshes.

Area-Weighted Averaging (Coarsening)




Figure 11: Area-weighted averaging for mesh coarsening: all source face centers (blue) within the search radius of each target point (orange) are averaged with equal weights

2D Projection by Box Face The interpolation is performed on 2D point clouds, as illustrated in fig. 12. Points are grouped by which face of the extraction bounding box they belong to (6 planar surfaces: $\pm X$, $\pm Y$, $\pm Z$), then projected to 2D by dropping the constant coordinate. This exploits the box geometry while handling the irregular face distribution that arises from extracting a submesh with original cell shapes.




Figure 12: Points on each box face are projected to 2D for triangulation. The constant coordinate (perpendicular to the face) is dropped, enabling efficient 2D algorithms

6.5.2 Initial Field Interpolation

Initial fields are also interpolated when mesh resolution changes (see figs. 13 and 14):

- **Refinement:** Uses `mapFields` with cell-center interpolation (fig. 13)
- **Coarsening:** Uses volume-weighted averaging of source cells (fig. 14)

i Note

Spatial interpolation introduces smoothing, particularly for coarsening. For turbulent flows, this may affect small-scale structures. Consider the trade-off between computational cost and resolution fidelity.

What `spaceTimeWindowInitCase` creates:

mapFields Cell-Center Interpolation (Refinement)




Figure 13: Initial field refinement using `mapFields`: target cell values are interpolated from surrounding source cell centers using distance-weighted averaging

Volume-Weighted Averaging (Coarsening)




Figure 14: Initial field coarsening using volume-weighted averaging: source cell values within each target cell region are averaged weighted by their volumes

1. `system/controlDict` – With matching solver, `deltaT`, `adjustTimeStep` from extraction
 - `startTime` set to t_2 (third timestep) for cubic interpolation buffer
 - `endTime` set to t_{n-2} (third-to-last) for cubic interpolation buffer
2. `system/fvSchemes`, `system/fvSolution` – Copied from source case (with `pRefPoint` added)
3. `constant/` files – All physics properties copied (mandatory for fidelity)
4. Initial field files with appropriate BCs based on options and `boundaryData` availability

7 Data Storage Format

7.1 Directory Structure

The extraction creates the following directory structure:

```

1  outputDir/
2      constant/
3          polyMesh/           # Subset mesh
4              points
5              faces
6              owner
7              neighbour
8              boundary
9              extractionBox     # (parallel only)
10             boundaryData/
11                 oldInternalFaces/
12                     points      # Face centres
13                     extractionMetadata # Settings and timestep list
14                     0.0001/
15                         U          # or U.dvz, U.dvz.zstd, U.dvz.zstd.enc
16                         p          # pressure (for coupled BC)
17                         gradp       # pressure gradient (for coupled BC)
18                     0.0002/
19                         ...
20             0.0001/           # Initial fields (serial only)
21                 U
22                 p
23                 nut

```

Listing 18: Extracted data structure

7.2 Boundary Data Compression

The `writeFormat` parameter controls how boundary data files are written, with compression ratios summarized in table 8.

Table 8: Compression comparison

Format	Extension	Typical Size	Notes
ASCII	(none)	100%	Human-readable
Binary	(none)	~50%	OpenFOAM native
ASCII + gzip	.gz	~10%	Compressed
Binary + gzip	.gz	~8%	Compressed
deltaVarint	.dvz	~2.7%	High compression, self-contained
deltaVarint + zstd	.dvz.zstd	~0.4%	Integrated zstd (auto if libzstd)
dvzt	.dvzt	~2.4%	Best compression, recommended
dvzt + zstd	.dvzt.zstd	~0.3%	Best compression + integrated zstd

7.2.1 Delta-Varint Codec (DVZ)

Specialized codec optimized for CFD boundary data (see figs. 15 and 16):

1. **Component-major ordering:** Groups similar values (all U_x , then U_y , then U_z)
2. **Spatial delta encoding:** Stores differences between consecutive face values within the same timestep
3. **Quantization:** Rounds to configurable precision
4. **Varint encoding:** Variable-length integer encoding
5. **Zigzag encoding:** Efficient signed integer representation

DVZ Encoding Pipeline




Figure 15: DVZ encoding pipeline: data is reordered by component for better locality, then delta-encoded spatially. Quantization, zigzag encoding (for signed integers), and variable-length integer encoding produce compact output




Figure 16: Spatial delta encoding: instead of storing absolute values, DVZ stores the first value and differences between consecutive faces. Smooth CFD fields have small deltas, enabling high compression

i Note

Each DVZ timestep file is completely self-contained. Delta encoding is purely spatial (between consecutive faces in a single field), not temporal. Any timestep can be read independently without access to other timesteps.

7.2.2 DVZ Mathematical Formulation

Component-Major Reordering For a vector field \mathbf{U} with N faces, the raw data layout is face-major (eq. (16)):

$$\text{Raw} : [U_{x,0}, U_{y,0}, U_{z,0}, U_{x,1}, U_{y,1}, U_{z,1}, \dots, U_{x,N-1}, U_{y,N-1}, U_{z,N-1}] \quad (16)$$

DVZ reorders to component-major for better compression (eq. (17)):

$$\text{Reordered} : \underbrace{[U_{x,0}, U_{x,1}, \dots, U_{x,N-1}]}_{\text{all } x}, \underbrace{[U_{y,0}, \dots, U_{y,N-1}]}_{\text{all } y}, \underbrace{[U_{z,0}, \dots, U_{z,N-1}]}_{\text{all } z} \quad (17)$$

Spatial Delta Encoding For each component sequence $\{v_0, v_1, \dots, v_{N-1}\}$, store deltas as shown in eq. (18):

$$\delta_0 = v_0, \quad \delta_i = v_i - v_{i-1} \quad \text{for } i = 1, \dots, N-1 \quad (18)$$

For smooth CFD fields, consecutive face values are similar, so $|\delta_i| \ll |v_i|$.

Quantization Convert floating-point deltas to integers with configurable precision p (decimal digits) using eq. (19):

$$q_i = \text{round}(\delta_i \times 10^p) \quad (19)$$

The precision parameter controls the trade-off between compression ratio and accuracy. With $p = 6$ (default), values are accurate to $\sim 10^{-6}$ relative precision.

Zigzag Encoding Convert signed integers to unsigned for efficient varint encoding (eq. (20)):

$$z_i = \begin{cases} 2q_i & \text{if } q_i \geq 0 \\ -2q_i - 1 & \text{if } q_i < 0 \end{cases} = (q_i \ll 1) \oplus (q_i \gg 31) \quad (20)$$

This maps small-magnitude signed integers to small unsigned integers: $0 \mapsto 0$, $-1 \mapsto 1$, $1 \mapsto 2$, $-2 \mapsto 3$, etc.

Variable-Length Integer Encoding Encode unsigned integers using 7 bits per byte, with the high bit indicating continuation (eq. (21)):

$$\text{bytes}(z) = \left\lceil \frac{\lfloor \log_2(z) \rfloor + 1}{7} \right\rceil \quad (21)$$

Small values ($z < 128$) use 1 byte; larger values use more bytes. Since smooth CFD fields produce small deltas, most values compress to 1–2 bytes instead of 4–8 bytes for raw floats/doubles.

7.2.3 Delta-Varint-Temporal Codec (DVZT)

Enhanced codec that exploits both spatial and temporal correlation for better compression (see figs. 17 to 19):

1. **Keyframes** (every N timesteps): Self-contained, same as DVZ (fig. 17)
2. **Delta frames**: Hybrid spatial-temporal prediction (fig. 18)
 - Uses weighted prediction: $\hat{v} = 0.3 \cdot v_{\text{spatial}} + 0.7 \cdot v_{\text{temporal}}$
 - Encodes residuals (actual - predicted) instead of raw spatial deltas
 - Typically ~10% smaller than DVZ for delta frames




Figure 17: DVZT frame structure: keyframes are self-contained and appear every N timesteps (configurable). Delta frames between keyframes use temporal prediction for better compression

```

1  writeFormat          dvzt;
2  deltaVarintPrecision 6;           // ~1e-6 relative precision
3  dvztKeyframeInterval 20;         // Keyframe every 20 timesteps (default)

```

Listing 19: DVZT configuration

7.2.4 DVZT Mathematical Formulation

Frame Types DVZT distinguishes between keyframes and delta frames based on the timestep index k (eq. (22)):

$$\text{Frame type}(k) = \begin{cases} \text{Keyframe} & \text{if } k \bmod K = 0 \\ \text{Delta frame} & \text{otherwise} \end{cases} \quad (22)$$

where K is the keyframe interval (default $K = 20$).




Figure 18: DVZT hybrid prediction: each value is predicted using 30% spatial neighbor (previous face at same timestep) and 70% temporal neighbor (same face at previous timestep). Only the small residual is stored




Figure 19: DVZT workflow: extraction writes compact .dvzt[zstd] files. Before reconstruction, spaceTimeWindowInitCase converts them to .dvz[zstd] format (required for random timestep access at runtime). The zstd compression layer is preserved through the conversion

Keyframe Encoding Keyframes use pure DVZ encoding (spatial delta only), as shown in eq. (23):

$$\delta_i^{(k)} = v_i^{(k)} - v_{i-1}^{(k)} \quad (\text{spatial delta}) \quad (23)$$

Delta Frame Hybrid Prediction Delta frames use a weighted combination of spatial and temporal neighbors (eq. (24)):

$$\hat{v}_i^{(k)} = \alpha \cdot v_{i-1}^{(k)} + (1 - \alpha) \cdot v_i^{(k-1)} \quad (24)$$

where $\alpha = 0.3$ (30% spatial weight, 70% temporal weight).

The residual to be encoded is given by eq. (25):

$$r_i^{(k)} = v_i^{(k)} - \hat{v}_i^{(k)} \quad (25)$$

For slowly varying flows with small Δt , consecutive timesteps are nearly identical, so $v_i^{(k)} \approx v_i^{(k-1)}$, making the temporal prediction highly accurate and $|r_i^{(k)}| \ll |v_i^{(k)} - v_{i-1}^{(k)}|$.

Compression Ratio Analysis Let σ_s^2 be the variance of spatial deltas and σ_t^2 the variance of temporal deltas. The hybrid prediction residual variance is approximately (eq. (26)):

$$\sigma_r^2 \approx \alpha^2 \sigma_s^2 + (1 - \alpha)^2 \sigma_t^2 \quad (26)$$

For small timesteps where $\sigma_t^2 \ll \sigma_s^2$ (eq. (27)):

$$\frac{\sigma_r}{\sigma_s} \approx \alpha = 0.3 \quad (27)$$

This explains the ~70% reduction in residual magnitude, which translates to significant compression improvement through varint encoding.

Decoding Dependency Delta frames depend on the previous frame for reconstruction (eq. (28)):

$$v_i^{(k)} = r_i^{(k)} + \alpha \cdot v_{i-1}^{(k)} + (1 - \alpha) \cdot v_i^{(k-1)} \quad (28)$$

This creates a dependency chain from each keyframe (eq. (29)):

$$\text{Keyframe } k_0 \rightarrow \text{Frame } k_0 + 1 \rightarrow \dots \rightarrow \text{Frame } k_0 + K - 1 \quad (29)$$

Random access requires decoding from the nearest preceding keyframe. For this reason, the initialization utility converts DVZT to DVZ before runtime.

DVZT Workflow:

During extraction, DVZT writes smaller .dvzt files (or .dvzt.zstd with integrated zstd). During case initialization, the utility automatically converts DVZT to DVZ format (required because delta frames need sequential processing). The zstd layer is preserved through the conversion. The resulting .dvz (or .dvz.zstd) files are read by the spaceTimeWindow BC at runtime.

Extraction → .dvzt[.zstd] files → spaceTimeWindowInitCase → .dvz[.zstd] files → Runtime

When to use DVZT:

- Long simulations with many timesteps (storage savings accumulate)
- Network/disk bandwidth constraints during extraction
- Small timesteps ($\Delta t = 10^{-5}$ or 10^{-6}) where temporal correlation is very strong
- DNS or acoustic simulations requiring fine temporal resolution

Compression vs Timestep Size:

DVZT benefits increase dramatically with smaller timesteps because consecutive values become nearly identical (see table 9):

Table 9: DVZT compression improvement by timestep size

Δt	Temporal Correlation	DVZT vs DVZ Savings
10^{-4}	Moderate	~10% smaller
10^{-5}	Strong	~20–30% smaller
10^{-6}	Very strong	~30–50% smaller

When to use DVZ:

- Simpler workflow (no conversion step)
- When random access to individual timesteps is needed during extraction
- Shorter simulations where DVZT overhead isn't worth it
- Large timesteps where temporal correlation is weak

7.2.5 Integrated Zstd Compression

When built with libzstd support (auto-detected by `./Allwmake`), DVZ and DVZT data is automatically wrapped with zstd level-3 compression before writing to disk. This is applied transparently as a compression layer in the encoding pipeline:

codec encode → zstd compress → [optional encrypt] → write to disk

Extension stacking:

- `fieldName.dvz.zstd` – DVZ with integrated zstd
- `fieldName.dvzt.zstd` – DVZT with integrated zstd
- `fieldName.dvz.zstd.enc` – DVZ + zstd + encryption
- `fieldName.dvzt.zstd.enc` – DVZT + zstd + encryption

The zstd layer provides an additional $\sim 10\times$ reduction on top of the already-compact varint encoding, bringing typical file sizes to $\sim 0.3\text{--}0.4\%$ of uncompressed ASCII. Decompression is transparent at runtime—the boundary conditions automatically detect and decompress `.zstd` files.

i Note

No configuration is needed: zstd compression is applied automatically when the library is built with libzstd. Files without the .zstd extension (from older builds or builds without libzstd) remain fully readable for backward compatibility.

The DVZT workflow with integrated zstd becomes:

.dvzt.zstd files → spaceTimeWindowInitCase → .dvz.zstd files → Runtime

7.2.6 External Compression Benchmark

i Note

With integrated zstd compression (built with libzstd), boundary data is already compressed on the fly with zstd -3. The benchmarks below show what is achievable with *external* tools on *uncompressed* DVZ/DVZT files. For archival of already zstd-compressed files, use 7z LZMA2 for additional gains.

DVZ and DVZT files can be further compressed using external tools for archival or transfer. The following benchmarks compare various compression algorithms on real boundary data files.

Test Environment:

- **CPU:** Intel Core i7-4790 @ 3.60 GHz (4 cores, 8 threads, Haswell microarchitecture)
- **L3 Cache:** 8 MB
- **RAM:** DDR3-1600

Test Data: 618 DVZ files totaling 13 MB (spatial delta-varint encoded), and 938 DVZT files totaling 30 MB (temporal delta-varint encoded).

The compression results are presented in table 10 for DVZ files and table 11 for DVZT files.

Table 10: External compression results for DVZ files (13 MB original)

Method	Size	Ratio	Speed	Notes
7z lzma2 -mx9	769 KB	6.01%	11.2 MB/s	Best ratio
xz -6	769 KB	6.01%	10.2 MB/s	
zstd -ultra -22	964 KB	7.53%	4.7 MB/s	Very slow
zstd -19	965 KB	7.54%	15.3 MB/s	
rar -m5	1017 KB	7.95%	26.9 MB/s	
7z ppmd -mx9	1.4 MB	11.02%	10.7 MB/s	
zstd -9	1.7 MB	13.17%	109.5 MB/s	
zstd -3	1.8 MB	13.65%	309.0 MB/s	Best balance
zstd -1	1.9 MB	14.43%	533.7 MB/s	
bzip2 -9	2.0 MB	15.28%	13.5 MB/s	
gzip -6	2.1 MB	16.39%	94.2 MB/s	
lz4	2.3 MB	17.80%	635.2 MB/s	Fastest

Key findings:

Table 11: External compression results for DVZT files (30 MB original)

Method	Size	Ratio	Speed	Notes
7z lzma2 -mx9	1.9 MB	6.31%	10.1 MB/s	Best ratio
xz -6	1.9 MB	6.32%	10.5 MB/s	
zstd -ultra -22	2.6 MB	8.48%	2.4 MB/s	Very slow
zstd -19	2.6 MB	8.50%	10.4 MB/s	
rar -m5	2.8 MB	9.28%	28.0 MB/s	
zstd -9	2.8 MB	9.27%	120.6 MB/s	
7z ppmf -mx9	2.8 MB	9.37%	12.2 MB/s	
zstd -3	3.1 MB	10.10%	396.1 MB/s	Best balance
zstd -1	3.1 MB	10.25%	614.0 MB/s	
bzip2 -9	3.5 MB	11.52%	13.7 MB/s	
gzip -6	3.8 MB	12.71%	104.3 MB/s	
lz4	4.1 MB	13.56%	763.0 MB/s	Fastest

- **Best compression ratio:** 7z LZMA2 and xz achieve ~6% (94% reduction)
- **Best speed/ratio balance:** zstd -3 at 10–14% ratio with 300–400 MB/s throughput
- **zstd -3 is 40× faster than xz** with only 60% more space
- **zstd -ultra -22 provides no benefit** over zstd -19 for this data type
- **bzip2 and PPMd perform poorly** for CFD boundary data

Recommendations:

- **Runtime/on-the-fly:** Use integrated zstd (build with libzstd, automatic)
- **Archive/transfer:** 7z lzma2 -mx9 or xz -6 (~6% ratio, 10 MB/s)
- **Real-time streaming:** lz4 (~14–18% ratio, 600–800 MB/s)

```

1 # Archive with best compression (7z)
2 cd subset-case/constant/boundaryData/oldInternalFaces
3 7z a -m0=lzma2 -mx=9 ../boundaryData.7z */*.dvz.zstd */*.dvz */*.dvzt.zstd
   */*.dvzt
4
5 # Or with tar (files already zstd-compressed if built with libzstd)
6 tar -cf ../boundaryData.tar */*.dvz.zstd

```

Listing 20: Archival compression example

7.3 Extraction Metadata

The extractionMetadata file contains:

```

1 {
2     openfoamVersion      "v2512";
3     openfoamApi          2512;

```

```

4     solver          "pimpleFoam";
5     deltaT         1e-04;
6     adjustTimeStep false;
7     timePrecision   6;
8     extractionStartTime 0.0001;
9     boxMin          (0.05 -0.25 0.01);
10    boxMax          (0.90 0.25 0.38);
11    nGlobalFaces    12345;
12    timesteps       (0.0001 0.0002 0.0003 ...);
13 }

```

Listing 21: extractionMetadata contents

8 Parallel Execution

Both the source simulation (extraction phase) and the reconstruction simulation can run in parallel. This enables efficient use of HPC resources for both phases of the workflow.

8.1 Parallel Extraction

The extraction function object fully supports parallel execution:

- Extraction box can span multiple processor domains
- Boundary data is gathered from all processors
- Master processor writes combined data files
- Processor boundary faces are handled automatically

```

1 # Decompose the case
2 decomposePar
3
4 # Run parallel extraction
5 mpirun -np 8 pimpleFoam -parallel
6
7 # Automatic field writes occur at:
8 #   t_0 - extraction start (for lookback)
9 #   t_1 - linear interpolation start
10 #  t_2 - cubic interpolation start

```

Listing 22: Parallel extraction

8.2 Field Reconstruction

After parallel extraction, reconstruct fields before case initialization:

```

1 # For cubic interpolation (recommended)
2 reconstructPar -time <t_2>
3
4 # For linear interpolation

```

```
5  reconstructPar -time <t_1>
```

Listing 23: Field reconstruction

⚠ Warning

The subset mesh created by `spaceTimeWindowInitCase` uses cell ordering from the reconstructed (serial) source mesh. Running `reconstructPar` before case initialization is mandatory for parallel extractions.

8.3 Parallel Reconstruction

The reconstruction simulation can also run in parallel, independently of how the extraction was performed:

```
1  cd subset-case
2  decomposePar
3  mpirun -np 4 pimpleFoam -parallel
4  reconstructPar
```

Listing 24: Parallel reconstruction

✓ Tip

The reconstruction case is typically much smaller than the source case, so fewer processors may be needed. The `spaceTimeWindow` boundary conditions work identically in serial and parallel modes.

9 Mass Conservation

9.1 The Mass Imbalance Problem

Face interpolation during extraction can introduce small mass flux imbalances:

$$\text{imbalance} = \sum_f \mathbf{U}_f \cdot \mathbf{S}_f \neq 0 \quad (30)$$

When all boundary patches have `fixesValue=true`, `adjustPhi()` cannot correct this imbalance, potentially causing pressure solver issues. Additionally, with all-Dirichlet BCs, the solver has no way to relieve pressure buildup.

9.2 Solutions

9.2.1 Use Pressure-Coupled BC (Recommended)

The `spaceTimeWindowCoupledPressure` BC handles mass conservation while accurately reproducing the original pressure field:

- Uses mixed Dirichlet/Neumann based on local flux
- Two modes: `fluxMagnitude` (default, Dirichlet at stagnant faces) or `flowDirection` (Dirichlet at inflow)

- Neumann faces allow natural pressure adjustment for continuity compatibility
- Requires gradient extraction during source simulation

See section 4.1 for detailed configuration.

9.2.2 Use -inletOutletBC (for quick tests)

The flux-based inlet-outlet BC naturally handles mass conservation:

- Outflow faces use zeroGradient, allowing natural outflow
- No artificial mass imbalance from prescribed outflow velocities
- Works without any special mass correction

```
1 spaceTimeWindowInitCase -sourceCase ../source -inletOutletBC
```

Listing 25: Using inlet-outlet BC for mass conservation

9.2.3 Use -correctMassFlux

Applies least-squares correction to boundaryData to ensure exact mass conservation:

$$\mathbf{U}_{\text{corrected}} = \mathbf{U} - \frac{\text{imbalance}}{\sum_f |\mathbf{S}_f|} \cdot \hat{\mathbf{n}} \quad (31)$$

where $\hat{\mathbf{n}} = \mathbf{S}_f / |\mathbf{S}_f|$ is the face unit normal.

This minimizes $\|\mathbf{U}_{\text{corrected}} - \mathbf{U}\|^2$ subject to:

$$\sum_f \mathbf{U}_{\text{corrected}} \cdot \mathbf{S}_f = 0 \quad (32)$$

```
1 spaceTimeWindowInitCase -sourceCase ../source -inletOutletBC -correctMassFlux
```

Listing 26: Using mass flux correction

9.2.4 Use -outletDirection with fixesValue true

Creates an outlet patch where mass imbalance can escape:

- oldInternalFaces uses fixesValue true (values not modified by adjustPhi)
- outlet uses inletOutlet BC (allows adjustPhi correction)

9.3 The fixesValue Option

The fixesValue parameter controls `adjustPhi()` behavior:

- `fixesValue = true`: Patch excluded from flux correction (preserves exact values)
- `fixesValue = false`: Patch included in flux correction (allows modification)

9.4 Outlet Patch for Pressure Relief

The `-outletDirection` option creates an outlet patch (see fig. 20):

```

1 spaceTimeWindowInitCase -sourceCase ./source \
2   -outletDirection "(1 0 0)" \
3   -outletFraction 0.1

```

Listing 27: Creating outlet patch




Figure 20: Outlet patch creation with `-outletDirection "(1 0 0)"`

Outlet faces are selected based on position (furthest along the outlet direction), not face orientation.

9.5 Recommended Configurations

The recommended mass conservation strategies are summarized in table 12.

Table 12: Mass conservation strategies

Configuration	Description
<code>-inletOutletBC</code>	Recommended for unsteady turbulent flows (vortex shedding, etc.). Natural mass balance through flux-based switching.
<code>-inletOutletBC -correctMassFlux</code>	Unsteady turbulent + extra safety. Ensures exact mass balance with inlet-outlet switching.
<code>-outletDirection "(1 0 0)"</code>	Steady-mean flow direction. Use when outlet location is known and flow direction is consistent.
<code>-correctMassFlux (no outlet)</code>	Maximum fidelity. All faces prescribed with exact mass balance.

10 Solver Settings and Relaxation

10.1 Time Stepping

For accurate reconstruction, use identical time stepping:

```

1  deltaT           1e-04;      // Must match extraction
2  adjustTimeStep no;          // Fixed timestep recommended
3
4  // If adaptive timestep is required:
5  adjustTimeStep yes;
6  maxCo            0.5;
7  maxDeltaT        1e-03;

```

Listing 28: Recommended controlDict settings**△ Warning**

The reconstruction must use identical `deltaT` and `adjustTimeStep` settings as the extraction. Mismatches cause fatal errors unless `allowTimeInterpolation=true`.

10.2 PIMPLE Settings

For incompressible flows:

```

1  PIMPLE
2  {
3    nOuterCorrectors   2;
4    nCorrectors        2;
5    nNonOrthogonalCorrectors 1;
6
7    // Pressure reference (set by spaceTimeWindowInitCase)
8    pRefPoint         (0.5 0 0.2);
9    pRefValue         0;
10 }

```

Listing 29: Recommended PIMPLE settings

10.3 Relaxation Factors

✓ Tip

For reconstruction simulations, relaxation is typically **not needed** because:

- Boundary conditions are prescribed (not coupled)
- Initial conditions come from the source simulation
- Flow field evolves smoothly from physical initial state

If stability issues occur, try:

```

1  relaxationFactors
2  {
3    fields
4    {
5      p             0.7;

```

```

6      }
7      equations
8      {
9          U           0.7;
10         "(k|epsilon|omega|nuTilda)" 0.7;
11     }
12 }
```

Listing 30: Optional relaxation

10.4 Linear Solver Settings

Standard settings work well:

```

1 solvers
2 {
3     p
4     {
5         solver      GAMG;
6         smoother   GaussSeidel;
7         tolerance   1e-06;
8         relTol     0.01;
9     }
10
11    U
12    {
13        solver      PBiCGStab;
14        preconditioner DILU;
15        tolerance   1e-06;
16        relTol     0.1;
17    }
18 }
```

Listing 31: Linear solver settings

11 Time Interpolation

The boundary conditions support three time interpolation modes, summarized in Table 13 and illustrated in fig. 21. The comparison between linear and cubic interpolation quality is shown in fig. 22.

Table 13: Time interpolation modes

Mode	Start Time	Timesteps Used	Use Case
none (exact)	t_0	1 (exact match)	“Bit”-reproducible results
linear	t_1	2 (bracketing)	Simple, smoothly varying flows
cubic	t_2	4 (Catmull-Rom)	Unsteady turbulent flows (recommended)

`spaceTimeWindowInitCase` automatically sets `startTime` and `endTime` to ensure sufficient buffer timesteps for the selected interpolation scheme (see fig. 23 for buffer requirements).




Figure 21: Time interpolation modes: exact matching uses only one timestep, linear interpolation uses two bracketing timesteps, and cubic (Catmull-Rom) uses four timesteps for smooth C^1 continuous interpolation




Figure 22: Comparison of linear and cubic interpolation: linear interpolation creates C^0 continuous curves with sharp corners at data points, while Catmull-Rom cubic splines provide C^1 continuity with smooth transitions




Figure 23: Time buffer requirements for each interpolation mode: exact matching can use the full range, linear needs one buffer timestep at the start, and cubic needs two buffer timesteps at both start and end

11.1 Exact Timestep Matching (Default)

By default, the boundary condition requires exact timestep matching:

- Simulation time must match available sample times (within 1% of `deltaT`)
- Fatal error if no matching timestep found
- Ensures “bit”-reproducible results (highest fidelity)

11.2 Linear Interpolation

For cases where exact matching is not possible:

```

1 oldInternalFaces
2 {
3     type spaceTimeWindow;
4     allowTimeInterpolation true;
5     timeInterpolationScheme linear;
6     // ...
7 }
```

Listing 32: Linear interpolation

Uses two bracketing timesteps (eq. (33)):

$$\text{value} = (1 - \alpha) \cdot v_i + \alpha \cdot v_{i+1} \quad \text{where} \quad \alpha = \frac{t - t_i}{t_{i+1} - t_i} \quad (33)$$

The interpolation parameter $\alpha \in [0, 1]$ represents the normalized position between the two samples. This is equivalent to first-order polynomial interpolation (eq. (34)):

$$\phi(t) = \phi_i + \frac{\phi_{i+1} - \phi_i}{t_{i+1} - t_i}(t - t_i) \quad (34)$$

Linear interpolation provides C^0 continuity (continuous values but discontinuous derivatives at sample points).

11.3 Cubic Spline Interpolation

For smoother results with adaptive timestepping:

```

1 oldInternalFaces
2 {
3     type spaceTimeWindow;
4     allowTimeInterpolation true;
5     timeInterpolationScheme cubic;
6     // ...
7 }
```

Listing 33: Cubic interpolation

Uses centripetal Catmull-Rom spline with four timesteps ($t_{i-1}, t_i, t_{i+1}, t_{i+2}$):

- Handles non-uniform time spacing correctly
- Provides C^1 continuity
- No overshoots or cusps
- Essential for `adjustTimeStep=yes`

11.3.1 Catmull-Rom Spline Formulation

The centripetal Catmull-Rom spline passes through all control points and provides C^1 continuity. Given four consecutive sample values $\phi_0, \phi_1, \phi_2, \phi_3$ at times t_0, t_1, t_2, t_3 , the interpolated value for $t \in [t_1, t_2]$ is computed as follows.

First, compute the centripetal parameterization distances (eq. (35)):

$$d_j = |t_{j+1} - t_j|^{0.5}, \quad j = 0, 1, 2 \quad (35)$$

The normalized parameter within the segment $[t_1, t_2]$ is (eq. (36)):

$$u = \frac{t - t_1}{t_2 - t_1} \quad (36)$$

The spline is constructed using the Barry-Goldman algorithm. Define intermediate points (eq. (37)):

$$\begin{aligned} A_1 &= \frac{t_2 - t}{t_2 - t_0} \phi_0 + \frac{t - t_0}{t_2 - t_0} \phi_1 \\ A_2 &= \frac{t_2 - t}{t_2 - t_1} \phi_1 + \frac{t - t_1}{t_2 - t_1} \phi_2 \\ A_3 &= \frac{t_3 - t}{t_3 - t_1} \phi_2 + \frac{t - t_1}{t_3 - t_1} \phi_3 \end{aligned} \quad (37)$$

Then compute (eq. (38)):

$$\begin{aligned} B_1 &= \frac{t_2 - t}{t_2 - t_0} A_1 + \frac{t - t_0}{t_2 - t_0} A_2 \\ B_2 &= \frac{t_3 - t}{t_3 - t_1} A_2 + \frac{t - t_1}{t_3 - t_1} A_3 \end{aligned} \quad (38)$$

Finally, the interpolated value is (eq. (39)):

$$\phi(t) = \frac{t_2 - t}{t_2 - t_1} B_1 + \frac{t - t_1}{t_2 - t_1} B_2 \quad (39)$$

This formulation correctly handles non-uniform time spacing, which is essential when adaptive timestepping (`adjustTimeStep=yes`) produces variable timesteps.

11.4 Time Range Requirements

- **Linear:** Needs 2 buffer timesteps at start
- **Cubic:** Needs 2 buffer timesteps at start and end
- `spaceTimeWindowInitCase` automatically sets appropriate `startTime` and `endTime`

i Note

Extrapolation outside the extraction window is **never** allowed, regardless of interpolation settings.

12 Flux Reporting and Diagnostics

Enable flux reporting to monitor mass conservation:

```

1 oldInternalFaces
2 {
3     type          spaceTimeWindow;
4     reportFlux   true;
5     // ...
6 }
```

Listing 34: Enabling flux reporting

Output format (fields explained in table 14):

```

1 spaceTimeWindow flux [oldInternalFaces] t=0.001 thisPatch=1.234e-06 (in=-0.0523 out=0.0523) | MESH TOTAL=5.678e-05 (fixed=-0.0523
adjustable=0.0523) | adjustPhi will correct: -5.678e-05
```

Listing 35: Flux report output

Table 14: Flux report fields

Field	Description
thisPatch	Net flux through this spaceTimeWindow patch
in	Sum of inward fluxes
out	Sum of outward fluxes
MESH TOTAL	Total net flux across all boundary patches
fixed	Flux from patches with <code>fixesValue=true</code>
adjustable	Flux from patches with <code>fixesValue=false</code>
adjustPhi will correct	Correction to be distributed

13 Troubleshooting

13.1 Common Errors

13.1.1 Time Step Mismatch

```

1 FOAM FATAL ERROR:
2 Time step mismatch between extraction and reconstruction!
3 Extraction deltaT: 8.53e-04
4 Current deltaT:    1e-03

```

Listing 36: deltaT mismatch error

Solution: Set `deltaT` in `system/controlDict` to match extraction.

13.1.2 No Matching Timestep

```

1 FOAM FATAL ERROR:
2 No exact timestep match found for t = 0.00015
3 Available times: 0.0001, 0.0002, 0.0003

```

Listing 37: Missing timestep error

Solution: Enable time interpolation with `allowTimeInterpolation true`.

13.1.3 Pressure Solver Divergence

Symptoms: Pressure residuals increase, NaN values appear.

Solutions:

1. Use `-correctMassFlux` option
2. Add outlet patch with `-outletDirection`
3. Set `fixesValue false` on `oldInternalFaces`
4. Check `pRefPoint` is inside the domain

13.1.4 Encryption Errors

```

1 FOAM FATAL ERROR:
2 Failed to decrypt file: constant/boundaryData/.../U.dvz.zstd.enc

```

Listing 38: Decryption failure

Solutions:

1. Verify private key is correct
2. Ensure library was built with `FOAM_USE_SODIUM=1`
3. Check file is not corrupted

14 Acknowledgments

This library implements the space-time window reconstruction method developed in [1]. The example case uses mesh generation from the ERCOFTAC Classic Collection Database [2].

A Example Case: ERCOFTAC UFR2-02 Square Cylinder

The library includes a complete example based on the ERCOFTAC UFR2-02 benchmark case [2, 3] (Case 043 in the ERCOFTAC database): turbulent flow around a square cylinder at $Re = 21,400$. This case demonstrates vortex shedding and is ideal for testing space-time window reconstruction.

OPENFOAM® is a registered trademark of OpenCFD Limited.

A.1 Reference Data and Validation Resources

The ERCOFTAC (European Research Community on Flow, Turbulence and Combustion) Classic Collection Database provides extensive reference data for this test case, including:

- Experimental measurements from Lyn et al. [3] (laser-Doppler velocimetry)
- Reference numerical simulations from various research groups
- Mesh generation scripts and case setup files
- Detailed flow statistics and validation data

The database is accessible at:

<http://cfd.mace.manchester.ac.uk/ercoftac/>

An OpenFOAM-specific implementation guide is available at:

https://openfoamwiki.net/index.php?title=Benchmark_ercoftac_ufr2-02

i Note

If HTTPS access fails for the ERCOFTAC database, use HTTP (<http://>) instead. The ERCOFTAC server may not support secure connections.

A.2 Physical Background: Von Kármán Vortex Street

When fluid flows past a bluff body (such as a square cylinder), the flow separates at the sharp edges, creating alternating vortices that are shed from each side of the body. This phenomenon, known as the **von Kármán vortex street**, is characterized by:

- **Periodic vortex shedding:** Vortices detach alternately from upper and lower surfaces
- **Strouhal number:** $St = fD/U_\infty \approx 0.13$ for square cylinders, where f is the shedding frequency, D is the cylinder side length, and U_∞ is the freestream velocity
- **Turbulent wake:** At $Re = 21,400$, the wake is fully turbulent with complex three-dimensional structures
- **Unsteady forces:** Fluctuating lift and drag on the cylinder

The vortex street extends far downstream, making it an excellent test case for space-time window extraction—the extraction region can capture the wake dynamics while excluding the cylinder and inlet regions.

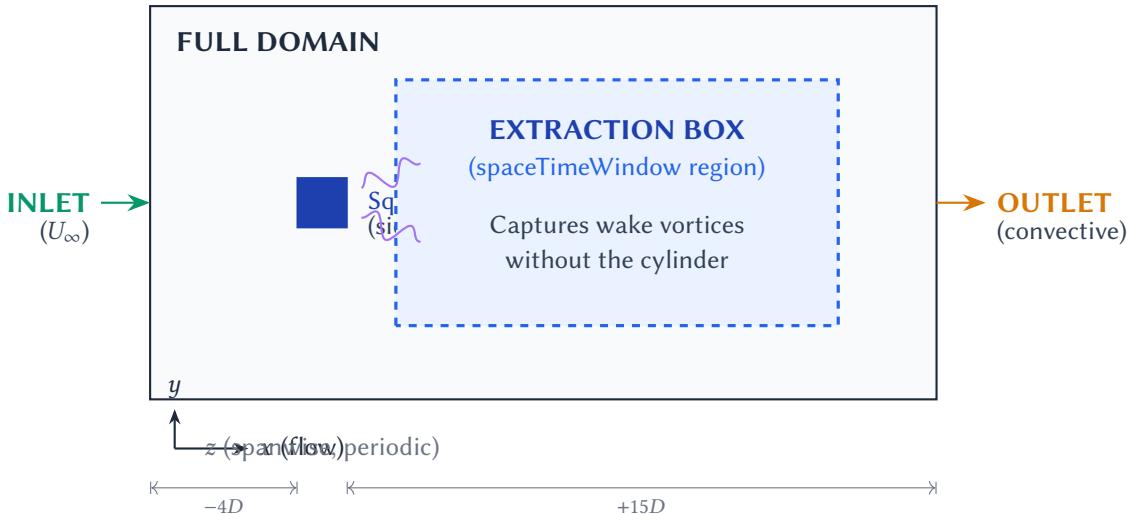


Figure 24: Computational domain for the square cylinder case showing the extraction box placement in the wake region

A.3 Domain and Mesh Configuration

The computational domain (fig. 24) extends from $-4D$ upstream to $+15D$ downstream of the cylinder, with lateral boundaries at $\pm 6.5D$. The spanwise extent is $4D$ with periodic boundary conditions.

A.4 Extraction Box Placement

The extraction box is positioned to capture the vortex street while avoiding:

- The cylinder itself (to avoid complex geometry in subset)
- The inlet region (where flow is uniform and uninteresting)
- Domain boundaries (box must be fully internal)

Example extraction configuration for this case:

```

1 extractSubset
2 {
3     type      spaceTimeWindowExtract;
4     libs      (spaceTimeWindow);
5
6     // Box starts 0.5D downstream of cylinder, extends to 9D
7     // Lateral extent captures the wake width
8     // Full spanwise extent (periodic direction)
9     box       ((0.05 -0.25 0.01) (0.90 0.25 0.38));
10
11    outputDir   ".../cylinder-subset";
12    fields      (U p nut);
13    writeFormat  deltaVarint;
14
15    writeControl timeStep;
16    writeInterval 1;

```

17 }

Listing 39: Extraction configuration for square cylinder

A.5 Physical Parameters

The physical parameters for this case are listed in table 15.

Table 15: Square cylinder case parameters

Parameter	Value	Description
D	0.04 m	Cylinder side length
U_∞	8.25 m/s	Freestream velocity
ν	1.5×10^{-5} m ² /s	Kinematic viscosity
Re	21,400	Reynolds number
St	≈ 0.13	Strouhal number
T_{shed}	≈ 0.037 s	Vortex shedding period

A.6 Running the Example

```

1 cd examples/ufr2-02
2
3 # Generate mesh and run full simulation
4 ./Allrun
5
6 # The extraction creates ../ufr2-02-subset with:
7 #   - Subset mesh in constant/polyMesh
8 #   - Boundary data in constant/boundaryData
9 #   - Initial fields
10
11 # Initialize and run reconstruction (recommended: inlet-outlet BC)
12 cd ..../ufr2-02-subset
13 spaceTimeWindowInitCase -sourceCase ../ufr2-02 -inletOutletBC
14
15 pimpleFoam
16
17 # Alternative: fixed outlet direction for steady-mean flows
18 # spaceTimeWindowInitCase -sourceCase ../ufr2-02 \
19 #   -outletDirection "(1 0 0)" \
20 #   -correctMassFlux

```

Listing 40: Running the example case

A.7 Validation

The reconstruction can be validated by comparing:

- Instantaneous velocity fields at matching timesteps
- Vortex shedding frequency (should match original)

- Mean velocity profiles in the wake
- Reynolds stress components

i Note

The reconstruction should exactly reproduce the original flow within the extraction region (to solver tolerance) when using exact timestep matching. Small differences may occur at the outlet boundary due to the different boundary condition type.

B Field Selection by Turbulence Model

When configuring the extraction, include all fields required by your turbulence model. Table 16 lists recommended fields for common models.

Table 16: Recommended fields by turbulence model

Turbulence Model	Recommended Fields
LES Smagorinsky	(U p nut)
LES dynamicKEqn	(U p nut k)
LES WALE	(U p nut)
RANS k- ε	(U p nut k epsilon)
RANS k- ω SST	(U p nut k omega)
Spalart-Allmaras	(U p nut nutilda)

References

- [1] Alin-Adrian Anton. *Space-Time Window Reconstruction in High-Performance Numeric Simulations: Application for CFD*. PhD thesis, Universitatea Politehnica Timișoara, Timișoara, Romania, November 2011. URL <https://dspace.upt.ro/jspui/handle/123456789/643>.
- [2] ERCOFTAC. Classic collection database: Case 043 – flow around a square cylinder. <http://cfd.mace.manchester.ac.uk/ercoftac/>, 2024. Mesh generation script by Niklas Nordin.
- [3] D. A. Lyn, S. Einav, W. Rodi, and J.-H. Park. A laser-Doppler velocimetry study of ensemble-averaged characteristics of the turbulent near wake of a square cylinder. *Journal of Fluid Mechanics*, 304:285–319, 1995. doi: 10.1017/S0022112095004435.